


Lab3_ForecastEvaluation&Combination_a24kimwu

2025-09-16

Load data & relevant packages

```
load("./grocery.Rdata")
```

```
plot(y)
```




```
n <- length(y)
n
## [1] 364
library(forecast)

## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo
library(tsutils)
```

Separate into in- and out-of-sample and data exploration


```
y.train <- head(y, 7*50)
y.test <- tail(y, 7*2)
```

```
cma <- cmav(y.train, outplot = TRUE)
```



```
seasplot(y.train)
```


**Seasonal plot (Detrended)
Seasonal (p-val: 0)**



```
## Results of statistical testing  
## Evidence of trend: TRUE  (pval: 0)  
## Evidence of seasonality: TRUE  (pval: 0)
```


```
seasplot(y.train, outplot = 2)
```

Seasonal boxplot (Detrended) Seasonal (p-val: 0)




```
## Results of statistical testing
## Evidence of trend: TRUE  (pval: 0)
## Evidence of seasonality: TRUE  (pval: 0)
{seasplot(y.train, outplot = 3)}
seasplot(head(y.train,10*7)) # First 10 weeks only
```

Seasonal plot (Detrended) Seasonal (p-val: 0)




```
## Results of statistical testing
## Evidence of trend: TRUE  (pval: 0.001)
## Evidence of seasonality: TRUE  (pval: 0)
seasplot(tail(y.train,10*7)) # Last 10 weeks only
```


Seasonal plot Seasonal (p-val: 0)



```
## Results of statistical testing
## Evidence of trend: FALSE (pval: 0.045)
## Evidence of seasonality: TRUE (pval: 0)
dc <- decomp(y.train, outplot=TRUE)
```



```
dc <- decomp(y.train, outplot=TRUE, type="pure.seasonal")
```



Forecasting

Selection of forecasts using information criteria

```
fit <- ets(y.train)
fit

## ETS(M,N,M)
##
## Call:
## ets(y = y.train)
##
## Smoothing parameters:
##   alpha = 0.1024
##   gamma = 0.3636
##
## Initial states:
##   l = 14657.0119
##   s = 1.2821 1.2869 0.993 0.8979 0.7854 0.8323
##                   0.9224
##
## sigma: 0.1747
##
##      AIC      AICc      BIC
## 7537.494 7538.142 7576.073

# Level model
fit1 <- ets(y.train,model="ANN")
# Seasonal model
fit2 <- ets(y.train,model="ANA")
# Linear trend model
fit3 <- ets(y.train,model="AAN",damped=FALSE)
# Damped trend model
fit4 <- ets(y.train,model="AAN",damped=TRUE)
# Trend seasonal model
fit5 <- ets(y.train,model="AAA",damped=FALSE)
# Damped trend seasonal model
fit6 <- ets(y.train,model="AAA",damped=TRUE)

aicc <- c(fit1$aicc,fit2$aicc,fit3$aicc,fit4$aicc,fit5$aicc,fit6$aicc)
# We name the aicc vectr to easily identify which is which
names(aicc) <- c("ANN", "ANA", "AAN", "AAdN", "AAA", "AAAdA")
aicc

##      ANN      ANA      AAN      AAdN      AAA      AAAdA
## 7727.092 7573.495 7732.440 7767.566 7578.661 7580.221

which.min(aicc)

## ANA
## 2
fit$aicc

## [1] 7538.142
fit2$aicc
```

```
## [1] 7573.495
```

Selection of forecasts using a validation set

```
y.ins <- head(y.train,48*7)
y.val <- tail(y.train,2*7)

h <- 7 #forecast 7 periods ahead

fit1v <- ets(y.ins,model="ANN")
fit2v <- ets(y.ins,model="ANA")
fit3v <- ets(y.ins,model="AAN",damped=FALSE)
fit4v <- ets(y.ins,model="AAN",damped=TRUE)
fit5v <- ets(y.ins,model="AAA",damped=FALSE)
fit6v <- ets(y.ins,model="AAA",damped=TRUE)
fit7v <- ets(y.ins,model="MNM")


frc1v<-forecast(fit1v,h=h)
frc2v<-forecast(fit2v,h=h)
frc3v<-forecast(fit3v,h=h)
frc4v<-forecast(fit4v,h=h)
frc5v<-forecast(fit5v,h=h)
frc6v<-forecast(fit6v,h=h)
frc7v<-forecast(fit7v,h=h)

#And the naive:
frc8v<-tail(y.ins,frequency(y.ins))[1:h] #that is copy the last season

#using the function frequency() to find how long is a season and copy the last h of those observations.


plot(frc1v)
```

Forecasts from ETS(A,N,N)



```
plot(frc6v)
```

Forecasts from ETS(A,Ad,A)



```
err1v<-mean(abs(y.val[1:h]-frc1v$mean))
```

#*y.val[1:h]* gives us the first 7 observations of *y.val*

#*frc1v\$mean* gives us the point forecasts of *frc1v*

#*mean(abs())* gives us the MAE.

```
err2v<-mean(abs(y.val[1:h]-frc2v$mean))
```

```
err3v<-mean(abs(y.val[1:h]-frc3v$mean))
```

```
err4v<-mean(abs(y.val[1:h]-frc4v$mean))
```

```
err5v<-mean(abs(y.val[1:h]-frc5v$mean))
```

```
err6v<-mean(abs(y.val[1:h]-frc6v$mean))
```

```
err7v<-mean(abs(y.val[1:h]-frc7v$mean))
```

#For the naive we just have the numeric values in *frc8v*, so we do not need the suffix \$mean
err8v<-mean(abs(y.val[1:h]-frc8v))

```
errv<-c(err1v,err2v,err3v,err4v,err5v,err6v,err7v,err8v)
```

```
names(errv)<-c("ANN", "ANA", "AAN", "AAdN", "AAA", "AAdA", "MNM", "Naive")
```

```
errv
```

```
##      ANN      ANA      AAN      AAdN      AAA      AAdA      MNM      Naive
## 2975.734 2822.253 3040.983 4786.245 2826.875 2822.616 2270.250 5367.263
```

```
which.min(errv)
```

```
## MNM
```

```
##    7
```

```
omax <- length(y.val)- h + 1
```

```
omax
```

```
## [1] 8
```

```

# This is what we will be running
models <- c("ANN", "ANA", "AAN", "AAN", "AAA", "AAA", "MNM", "Naive")
damped <- c(FALSE, FALSE, FALSE, TRUE, FALSE, TRUE, FALSE, FALSE)
# And this is where we will store things
# Forecast errors across forecast origins
err <- array(NA,c(omax,8)) # This has omакс rows and 8 columns, one for each different forecasting method
frcs <- array(NA,c(h,8))

for (o in 1:omax){
  print(o)
}

## [1] 1
## [1] 2
## [1] 3
## [1] 4
## [1] 5
## [1] 6
## [1] 7
## [1] 8

# For each forecast origin
for (o in 1:omax){
  # Split training set
  y.ins <- head(y.train,48*7-1+o) # As o increases, so will the in-sample.
  y.val <- tail(y.train,2*7-o+1) # As o increases, the validation will decrease.
  # Fit and forecast with all exponential smoothing models
  for (m in 1:7){
    fitTemp <- ets(y.ins,model=models[m],damped=damped[m])
    frcs[,m] <- forecast(fitTemp,h=h)$mean
    err[o,m] <- mean(abs(y.val[1:h]-frcs[,m]))
  }
  # Forecast using the seasonal naive
  # Remember we do not have a model for this
  frcs[,8] <- tail(y.ins,frequency(y.ins))[1:h]
  err[o,8] <- mean(abs(y.val[1:h]-frcs[,8]))
}

colnames(err) <- c("ANN", "ANA", "AAN", "AAdN", "AAA", "AAdA", "MNM", "Naive")
err

##          ANN      ANA      AAN     AAdN      AAA      AAdA      MNM      Naive
## [1,] 2975.734 2822.253 3040.983 4786.245 2826.875 2822.616 2270.250 5367.263
## [2,] 2603.560 2459.703 2647.831 5340.174 2461.724 2452.621 2172.020 5019.013
## [3,] 2262.265 2055.915 1945.982 3877.711 2061.953 2050.249 2248.446 4842.480
## [4,] 2434.633 1988.818 2449.286 2160.198 1997.578 1988.870 2488.761 4109.314
## [5,] 2530.054 1843.592 2574.834 5254.436 1856.715 1847.040 2251.791 3647.554
## [6,] 2472.729 1813.034 2568.853 4881.620 1815.201 1811.906 1781.343 3019.866
## [7,] 2816.176 1684.053 2903.166 2208.285 1684.137 1684.079 1606.699 2634.483
## [8,] 3148.847 1546.037 3233.429 2379.490 1546.174 1545.162 1568.200 2400.111

errMean <- colMeans(err)
errMean

##          ANN      ANA      AAN     AAdN      AAA      AAdA      MNM      Naive
## 2655.500 2026.676 2670.545 3861.020 2031.295 2025.318 2048.439 3880.011


```

```
which.min(errMean)
```

```
## AAdA
```

```
##     6
```

```
boxplot(err)
```



Out-of-sample evaluation

```
# What to run
modelsTest <- c("ANA", "MNM", "AAA", "Naive", "CombMean", "CombMedian")
dampedTest <- c(FALSE, FALSE, TRUE)

# Pre-allocate memory
omaxTest <- length(y.test) - h + 1
errTest <- array(NA, c(omaxTest, 6))
frcsTest <- array(NA, c(h, 6))

# For each forecast origin
for (o in 1:omaxTest){

  # Split training set
  y.trnTest <- head(y, 50*7-1+o) # As o increases, so will the in-sample.
  y.tstTest <- tail(y, 2*7-o+1) # As o increases, the test will decrease.

  # Fit and forecast will all exponential smoothing models
  for (m in 1:3){
    fitTemp <- ets(y.trnTest, model=modelsTest[m], damped=dampedTest[m])
    frcsTest[,m] <- forecast(fitTemp, h=h)$mean
    errTest[o,m] <- mean(abs(y.tstTest[1:h] - frcsTest[,m]))
  }

  # Forecast using the seasonal naive
  frcsTest[,4] <- tail(y.trnTest, frequency(y.trnTest))[1:h]
```

```

errTest[o,4] <- mean(abs(y.tstTest[1:h] - frcsTest[,4]))


# Combinations
# The function apply allows us to use any function we want on a matrix
frcsTest[,5] <- apply(frcsTest[,1:4], 1, mean)
# This reads, using array frcsTest[,1:4], i.e. all rows and the first 4 columns take the mean across a
errTest[o,5] <- mean(abs(y.tstTest[1:h] - frcsTest[,5]))

# And for the median:
frcsTest[,6] <- apply(frcsTest[,1:4], 1, median)
errTest[o,6] <- mean(abs(y.tstTest[1:h] - frcsTest[,6]))
}

# Assign names to errors
colnames(errTest) <- c("ANA", "MNM", "AAdA", "Naive", "Comb.Mean", "Comb.Median")

# Summarise and plot errors
boxplot(errTest)

```



```

errTestMean <- colMeans(errTest)
print(errTestMean)

##          ANA          MNM          AAdA          Naive      Comb.Mean  Comb.Median
## 1208.019  1362.662  1208.888  1691.712   1057.314   1193.481

which.min(errTestMean)

## Comb.Mean
##      5

```

Forecast combination with AIC weights

```

y.train <- window(AirPassengers, end=c(1959,12))
y.test <- window(AirPassengers, start=c(1960,1))

```

```

models <- c("ANN", "AAN", "MNM", "MAM")

fit <- list() # Here I will store models
frc <- array(NA, c(12,4), dimnames=list(NULL, models))# Here I will store forecasts

for (i in 1:4){
  # I ask ets() everytime to fit the model specified in the models variable
  fit[[i]] <- ets(y.train, model=models[i], damped=FALSE)
  # And then give me the forecasts
  frc[,i] <- forecast(fit[[i]], h=12)$mean
}

AIC <- unlist(lapply(fit, function(x){x$aic}))
AIC

## [1] 1558.920 1562.628 1297.518 1257.573

dAIC <- AIC - min(AIC)
dAIC <- exp(-0.5*dAIC)
waic <- dAIC/sum(dAIC)
waic


## [1] 3.659296e-66 5.730754e-67 2.118497e-09 1.000000e+00

round(waic, 4)

## [1] 0 0 0 1

plot(AirPassengers)

```



```

#Prepare variables and models
fit2<-list()
frc2<-array(NA,c(12,6))
models<-rep(c("AAA", "MAM", "MMM"), 2)
damped<-c(rep(FALSE, 3), rep(TRUE, 3))

```

```

#Fit models and generate forecasts
for(i in 1:6){
  fit2[[i]] <- ets(y.train, model=models[i], damped=damped[i])
  frc2[,i] <- forecast(fit2[[i]], h=12)$mean
}

#Extract AIC and calculate weights
AIC2<-unlist(lapply(fit2,function(x){x$aic}))
dAIC2<-AIC2-min(AIC2)
dAIC2<-exp(-0.5*dAIC2)
waic2<-dAIC2/sum(dAIC2)
round(waic2,4)

## [1] 0.0000 0.0005 0.0001 0.0000 0.3481 0.6513

# AIC weights
frcComb <- frc2 %*% cbind(waic2)
# Mean
frcComb <- cbind(frcComb, rowMeans(frc2))
# To calculate the median without loops we use apply()
frcComb <- cbind(frcComb, apply(frc2, 1, median))
# Selection
frcComb <- cbind(frcComb, frc2[,which.min(AIC2)])
colnames(frcComb) <- c("Comb.AIC", "Comb.Mean", "Comb.Median", "Selection")

err <- matrix(rep(y.test, 4), ncol=4) - frcComb
# The matrix(rep(y.tst, 4), ncol=4) creates a 4 column matrix with copies
# of the test set. Run it on its own to see the result.
MAE <- colMeans(abs(err))
round(MAE, 2)

##      Comb.AIC    Comb.Mean Comb.Median   Selection
##        22.03       20.74      20.99      21.64

```

Exercises

1. Grocery time series with increased validation & test set

```

y.train_1 <- head(y, 7*45)
y.test_1 <- tail(y, 7*7)

h_1 <- 7
y.val_1 <- tail(y.train_1, 5*7)

omax_1 <- length(y.val_1) - h_1 + 1
omax_1

## [1] 29

models_1 <- c("ANN", "ANA", "AAN", "AAN", "AAA", "AAA", "MNM", "Naive")
damped_1 <- c(FALSE, FALSE, FALSE, TRUE, FALSE, TRUE, FALSE, FALSE)
# And this is where we will store things
# Forecast errors across forecast origins
err_1 <- array(NA, c(omax_1, 8)) # This has omax rows and 8 columns, one for each different forecasting m
frcs_1 <- array(NA, c(h_1, 8))

```

```

for (o in 1:omax_1){
  # Split training set
  y.ins_1 <- head(y.train_1,40*7-1+o) # As o increases, so will the in-sample.
  y.val_1 <- tail(y.train_1,5*7-o+1) # As o increases, the validation will decrease.
  # Fit and forecast with all exponential smoothing models
  for (m in 1:7){
    fitTemp <- ets(y.ins_1,model=models_1[m],damped=damped_1[m])
    frcs_1[,m] <- forecast(fitTemp,h=h_1)$mean
    err_1[o,m] <- mean(abs(y.val_1[1:h_1]-frcs_1[,m]))
  }
  # Forecast using the seasonal naive
  # Remember we do not have a model for this
  frcs_1[,8] <- tail(y.ins_1,frequency(y.ins_1))[1:h_1]
  err_1[o,8] <- mean(abs(y.val_1[1:h_1]-frcs[,8]))
}

colnames(err_1) <- c("ANN", "ANA", "AAN", "AAdN", "AAA", "AAdA", "MNM", "Naive")
err_1

```

	ANN	ANA	AAN	AAdN	AAA	AAdA	MNM
[1,]	3076.380	866.2621	3170.189	3287.858	834.6769	867.8297	1171.3082
[2,]	2765.747	647.9016	2845.267	4435.224	600.8809	649.1710	1172.2796
[3,]	2560.926	559.0736	2593.698	4707.484	490.1267	560.8097	1297.2553
[4,]	2599.492	572.5887	2646.674	2743.252	530.7351	573.8516	1102.0152
[5,]	2892.196	666.3296	3008.532	3003.780	637.0696	666.0890	908.5401
[6,]	5119.186	782.7705	2908.734	5120.772	772.4585	780.1582	831.9916
[7,]	5883.344	1119.6824	2565.299	5885.173	1127.6279	1114.3946	1002.2140
[8,]	2447.036	1191.3955	2832.984	2446.975	1218.2989	1189.6227	1077.3430
[9,]	1844.779	1200.8504	2960.803	1844.662	1245.5068	1199.4369	1160.3067
[10,]	2405.196	1345.2460	3204.149	2405.311	1410.2037	1343.5889	1382.2688
[11,]	2237.229	1425.2719	3277.843	2236.743	1493.0938	1424.1347	1422.0604
[12,]	2350.886	1195.6472	2492.960	2350.831	1239.6615	1194.2941	1030.4943
[13,]	5113.489	1170.3009	3289.784	5114.129	1196.8910	1169.4577	1054.1952
[14,]	4523.109	850.5385	3577.045	4523.199	859.1260	850.0237	824.9574
[15,]	3516.249	832.6461	3621.869	2877.673	839.7686	832.3124	840.3402
[16,]	3108.702	846.3900	3643.322	3109.232	852.1939	846.0125	878.1179
[17,]	4224.938	643.8622	3482.535	4226.189	648.0600	643.8061	742.8407
[18,]	2826.906	656.3456	3510.362	2826.685	659.6553	656.2349	787.2232
[19,]	3327.163	841.9668	3748.059	3327.249	854.7893	841.4500	849.4168
[20,]	6469.714	1002.2144	3474.833	6471.660	779.7502	1000.8742	897.5182
[21,]	5918.778	1031.0938	3403.668	5920.413	1060.5062	1029.6500	871.2393
[22,]	3092.661	921.1452	3199.755	2302.603	948.0616	919.2962	712.5812
[23,]	2726.961	1005.4271	3295.102	2727.541	1025.2691	1003.7794	769.4352
[24,]	3169.724	1217.1184	2991.479	3170.551	1225.4160	1215.7441	1018.8763
[25,]	2504.924	1132.4904	2919.029	2505.595	1139.7473	1131.2056	926.1661
[26,]	3051.948	1614.3301	2800.099	3052.180	1609.9510	1613.5378	1488.2227
[27,]	3819.308	1578.6880	3407.908	3819.329	1578.4906	1578.5699	1676.3688
[28,]	5035.723	2482.9313	3327.827	5036.575	2481.6473	2482.7784	2535.6730
[29,]	3052.491	2762.3666	3706.984	3052.334	2763.8088	2761.9646	2714.8382
							Naive
[1,]	2368.146						
[2,]	1456.426						
[3,]	3044.619						
[4,]	4926.071						

```

## [5,] 5104.420
## [6,] 5019.570
## [7,] 4365.256
## [8,] 2615.930
## [9,] 1427.506
## [10,] 3372.479
## [11,] 4811.429
## [12,] 5067.996
## [13,] 4991.063
## [14,] 4582.120
## [15,] 2794.916
## [16,] 1240.113
## [17,] 3412.226
## [18,] 5544.926
## [19,] 5673.607
## [20,] 5004.361
## [21,] 4207.309
## [22,] 2395.834
## [23,] 1449.031
## [24,] 2588.604
## [25,] 4522.967
## [26,] 5421.170
## [27,] 5762.084
## [28,] 3145.049
## [29,] 1890.356

errMean_1 <- colMeans(err_1)
errMean_1


##      ANN      ANA      AAN      AAAdN      AAA      AAAdA      MNM      Naive
## 3505.696 1109.065 3169.200 3604.524 1107.706 1108.279 1142.969 3731.227

which.min(errMean_1)

## AAA
## 5

boxplot(err_1)

```



```

# What to run
modelsTest_1 <- c("ANA", "MNM", "AAA", "Naive", "CombMean", "CombMedian")
dampedTest_1 <- c(FALSE, FALSE, TRUE)

# Pre-allocate memory
omaxTest_1 <- length(y.test_1) - h_1 + 1
errTest_1 <- array(NA, c(omaxTest_1, 6))
frcsTest_1 <- array(NA, c(h_1, 6))

# For each forecast origin
for (o in 1:omaxTest_1){

  # Split training set
  y.trnTest_1 <- head(y, 45*7-1+o) # As o increases, so will the in-sample.
  y.tstTest_1 <- tail(y, 7*7-o+1) # As o increases, the test will decrease.

  # Fit and forecast will all exponential smoothing models
  for (m in 1:3){
    fitTemp <- ets(y.trnTest_1, model=modelsTest_1[m], damped=dampedTest_1[m])
    frcsTest_1[,m] <- forecast(fitTemp, h=h_1)$mean
    errTest_1[o,m] <- mean(abs(y.tstTest_1[1:h_1] - frcsTest_1[,m]))
  }

  # Forecast using the seasonal naive
  frcsTest_1[,4] <- tail(y.trnTest_1, frequency(y.trnTest_1))[1:h_1]
  errTest_1[o,4] <- mean(abs(y.tstTest_1[1:h_1] - frcsTest_1[,4]))

  # Combinations
  # The function apply allows us to use any function we want on a matrix
  frcsTest_1[,5] <- apply(frcsTest_1[,1:4], 1, mean)
  # This reads, using array frcsTest[,1:4], i.e. all rows and the first 4 columns take the mean across a
  errTest_1[o,5] <- mean(abs(y.tstTest_1[1:h_1] - frcsTest_1[,5]))
  # And for the median:
  frcsTest_1[,6] <- apply(frcsTest_1[,1:4], 1, median)
  errTest_1[o,6] <- mean(abs(y.tstTest_1[1:h_1] - frcsTest_1[,6]))
}

```


```

}

# Assign names to errors
colnames(errTest_1) <- c("ANA", "MNM", "AAdA", "Naive", "Comb.Mean", "Comb.Median")

# Summarise and plot errors
boxplot(errTest_1)

```



```

errTestMean_1 <- colMeans(errTest_1)
print(errTestMean_1)

##          ANA           MNM           AAdA           Naive       Comb.Mean   Comb.Median
## 2422.541  2512.475  2423.530  2948.417  2463.114  2443.858
which.min(errTestMean_1)

```

```

## ANA
## 1

```

Question: Do the results change?

Answer: Using a larger test & validation set leads to ANA being the model with the lowest error in the test data, rather than Comb.Mean with the smaller test & validation set and AAA being the model best performing in validation set rather than AAdA.

2. Air passengers with rolling origin

```

y <- AirPassengers
y.train_2 <- window(AirPassengers, end=c(1959,12))
y.test_2 <- window(AirPassengers, start=c(1960,1))

print(head(y, 12*11))

##      Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
## 1949 112 118 132 129 121 135 148 148 136 119 104 118
## 1950 115 126 141 135 125 149 170 170 158 133 114 140
## 1951 145 150 178 163 172 178 199 199 184 162 146 166

```

```

## 1952 171 180 193 181 183 218 230 242 209 191 172 194
## 1953 196 196 236 235 229 243 264 272 237 211 180 201
## 1954 204 188 235 227 234 264 302 293 259 229 203 229
## 1955 242 233 267 269 270 315 364 347 312 274 237 278
## 1956 284 277 317 313 318 374 413 405 355 306 271 306
## 1957 315 301 356 348 355 422 465 467 404 347 305 336
## 1958 340 318 362 348 363 435 491 505 404 359 310 337
## 1959 360 342 406 396 420 472 548 559 463 407 362 405
h_2 <- 8

# What to run
modelsTest_2 <- c("AAA", "MAM", "MMM", "AAA", "MAM", "MMM", "Comb.AIC", "Comb.Mean", "Comb.Median", "Select")
dampedTest_2 <- c(rep(FALSE, 3), rep(TRUE, 3))

# Pre-allocate memory
omaxTest_2 <- length(y.test_2) - h_2 + 1
errTest_2 <- array(NA, c(omaxTest_2, 10))
frcsTest_2 <- array(NA, c(h_2, 10))

# For each forecast origin
for (o in 1:omaxTest_2){

  # Split training set
  y.trnTest_2 <- head(y, 12*11 - 1+o) # As o increases, so will the in-sample
  y.tstTest_2 <- tail(y, 12 - o+1) # As o increases, the test will decrease.

  # Fit and forecast will all exponential smoothing models
  for (m in 1:6){
    fitTemp <- ets(y.trnTest_2, model=modelsTest_2[m], damped=dampedTest_2[m])
    frcsTest_2[,m] <- forecast(fitTemp, h=h_2)$mean
    errTest_2[o,m] <- mean(abs(y.tstTest_2[1:h_2] - frcsTest_2[,m]))
    AIC2 <- AIC(fitTemp)
    waic2 [m] <- exp(-0.5 * (AIC2 - min(AIC2))) / sum(exp(-0.5 * (AIC2 - min(AIC2))))
  }

  # Forecast using AIC weights
  frcsTest_2[,7] <- cbind(y.tstTest_2[1:4] %*% cbind(waic2[1:4]))
  errTest_2[o,7] <- mean(abs(y.tstTest_2[1:h_2] - frcsTest_2[,7]))

  # Forecast using Mean
  frcsTest_2[,8] <- cbind(apply(frcsTest_2[,1:4], 1, mean))
  errTest_2[o,8] <- mean(abs(y.tstTest_2[1:h_2] - frcsTest_2[,8]))


  # Forecast using median
  frcsTest_2[,9] <- cbind(apply(frcsTest_2[,1:4], 1, median))
  errTest_2[o,9] <- mean(abs(y.tstTest_2[1:h_2] - frcsTest_2[,9]))

  # Forecast using Selection
  frcsTest_2[,10] <- cbind(frcsTest_2[, which.min(AIC2)])
  errTest_2[o,10] <- mean(abs(y.tstTest_2[1:h_2] - frcsTest_2[,10]))
}

# Assign names to errors
colnames(errTest_2) <- c("AAA", "MAM", "MMM", "AAA", "MAM", "MMM", "Comb.AIC", "Comb.Mean", "Comb.Median")

```

```
# Summarise and plot errors
boxplot(errTest_2)
```



```
errTestMean_2 <- colMeans(errTest_2)
print(errTestMean_2)
```

```
##          AAA          MAM          MMM          AAA          MAM          MMM
##  52.62375  14.19167  16.47041  57.83518  37.79167  37.63183
##  Comb.AIC  Comb.Mean Comb.Median Selection
## 1426.05000 28.20096  29.38835  52.62375
```

```
which.min(errTestMean_2)
```

```
## MAM
## 2
```

Question: Do the results change?

Answer: Using the rolling origin for the AirPassenger dataset lead to MAM being the most accurate model, instead of Comb.Mean when using the normal evaluation.